

# Step 1 Data mining

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv('train.csv')
```

```
In [3]: df.head()
```

Out[3]:

|   | PassengerId | Survived | Pclass | Name  | Sex    | Age  | SibSp | Parch | Ticket           | Fare    |
|---|-------------|----------|--------|---|--------|------|-------|-------|------------------|---------|
| 0 | 1           | 0        | 3      | Braund, Mr. Owen Harris                           | male   | 22.0 | 1     | 0     | A/5 21171        | 7.2500  |
| 1 | 2           | 1        | 1      | Cumings, Mrs. John Bradley (Florence Briggs Th... | female | 38.0 | 1     | 0     | PC 17599         | 71.2833 |
| 2 | 3           | 1        | 3      | Heikkinen, Miss. Laina                            | female | 26.0 | 0     | 0     | STON/O2. 3101282 | 7.9250  |
| 3 | 4           | 1        | 1      | Futrelle, Mrs. Jacques Heath (Lily May Peel)      | female | 35.0 | 1     | 0     | 113803           | 53.1000 |
| 4 | 5           | 0        | 3      | Allen, Mr. William Henry                          | male   | 35.0 | 0     | 0     | 373450           | 8.0500  |

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived    891 non-null    int64
2   Pclass      891 non-null    int64
3   Name        891 non-null    object
4   Sex         891 non-null    object
5   Age        714 non-null    float64
6   SibSp       891 non-null    int64
7   Parch       891 non-null    int64
8   Ticket      891 non-null    object
9   Fare        891 non-null    float64
10  Cabin       204 non-null    object
11  Embarked    889 non-null    object
```

```
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.7+ KB
```

```
In [5]: df.shape
```

```
Out[5]: (891, 12)
```

```
In [6]: # drop useless column and the column missing many data  
df=df.drop(columns=['PassengerId','Name','Cabin','Ticket'])
```

```
In [7]: #combine sibling number with parents number to family number  
df['family']=df['SibSp']+df['Parch']
```

```
In [8]: df=df.drop(columns=['SibSp','Parch'])
```

```
In [9]: df['is_alone']=df['family'].apply(lambda x:1 if x==0 else 0)
```

```
In [10]: # trun sex to dummy variable  
df['is_female']= df['Sex'].apply(lambda x:1 if x =='female' else 0 )
```


```
In [11]: df=df.drop(columns=['Sex'])
```

```
In [12]: # check the distribute of age  
sns.distplot(df[df['Age'].notnull()][ 'Age' ])
```

/Users/jameswang/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

```
Out[12]: <AxesSubplot:xlabel='Age', ylabel='Density'>
```



```
In [13]: df['Age']=df['Age'].interpolate()
```

```
In [14]:
```

```
df['is_child']=df['Age'].apply(lambda x:1 if x<18 else 0)
```

```
In [15]: #fill the Embarked
df['Embarked'].value_counts()
```

```
Out[15]: S    644
         C    168
         Q     77
         Name: Embarked, dtype: int64
```

```
In [16]: df['Embarked']=df['Embarked'].fillna('S')
```

```
In [17]: df=pd.get_dummies(df,prefix='is')
```

```
In [18]: df.head()
```

```
Out[18]:
```

|   | Survived | Pclass | Age  | Fare    | family | is_alone | is_female | is_child | is_C | is_Q | is_S |
|---|----------|--------|------|---------|--------|----------|-----------|----------|------|------|------|
| 0 | 0        | 3      | 22.0 | 7.2500  | 1      | 0        | 0         | 0        | 0    | 0    | 1    |
| 1 | 1        | 1      | 38.0 | 71.2833 | 1      | 0        | 1         | 0        | 1    | 0    | 0    |
| 2 | 1        | 3      | 26.0 | 7.9250  | 0      | 1        | 1         | 0        | 0    | 0    | 1    |
| 3 | 1        | 1      | 35.0 | 53.1000 | 1      | 0        | 1         | 0        | 0    | 0    | 1    |
| 4 | 0        | 3      | 35.0 | 8.0500  | 0      | 1        | 0         | 0        | 0    | 0    | 1    |

```
In [19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Survived    891 non-null    int64
 1   Pclass      891 non-null    int64
 2   Age         891 non-null    float64
 3   Fare        891 non-null    float64
 4   family      891 non-null    int64
 5   is_alone    891 non-null    int64
 6   is_female   891 non-null    int64
 7   is_child    891 non-null    int64
 8   is_C        891 non-null    uint8
 9   is_Q        891 non-null    uint8
10   is_S        891 non-null    uint8
dtypes: float64(2), int64(6), uint8(3)
memory usage: 58.4 KB
```

```
In [20]: x , y = df.drop(columns=['Survived']), df['Survived']
```

## Step 2 Build model

```
In [29]: from sklearn.linear_model import LogisticRegression
         from sklearn.svm import SVC, LinearSVC
         from sklearn.tree import DecisionTreeClassifier
         from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
```

```
In [23]: train_x, test_x, train_y, test_y = train_test_split(x,y,train_size=0.8,
                                                         test_size=0.2,random_state=42)
```

## Logistic Regression

```
In [82]: model_LR = LogisticRegression()
model_LR.fit(train_x,train_y)
pred_y=model_LR.predict(test_x)
acc_log_reg=round(model_LR.score(test_x,test_y)*100,2)
print(str(acc_log_reg)+'percent')
```

83.24percent

/Users/jameswang/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear\_model/\_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n\_iter\_i = \_check\_optimize\_result(

```
In [53]: confusion_matrix(test_y, pred_y, labels=[1, 0])
```

```
Out[53]: array([[ 47,  20],
                [ 10, 102]])
```

```
In [41]: a=str((102+47)/(102+10+20+47))
print('Correct rate is '+ a +'percent')
```

Correct rate is 0.8324022346368715percent

## Decision Tree

```
In [75]: model_tree =DecisionTreeClassifier(max_depth=3, random_state=20)
model_tree.fit(train_x,train_y)
pred_y=model_tree.predict(test_x)
acc_log_reg=round(model_tree.score(test_x,test_y)*100,2)
print(str(acc_log_reg)+'percent')
```

81.56percent

```
In [76]: confusion_matrix(test_y, pred_y, labels=[1, 0])
```

```
Out[76]: array([[ 38,  29],
                [  4, 108]])
```

```
In [77]: a=str((38+108)/(38+29+4+108))
print('Correct rate is '+ a +'percent')
```

Correct rate is 0.8156424581005587percent

## SVC

```
In [58]: model_svc = SVC()
model_svc.fit(train_x, train_y)
pred_y=model_svc.predict(test_x)
acc_log_reg=round(model_svc.score(test_x, test_y)*100, 2)
print(str(acc_log_reg)+'percent')
```

68.16percent

```
In [59]: confusion_matrix(test_y, pred_y, labels=[1, 0])
```

```
Out[59]: array([[ 22,  45],
               [ 12, 100]])
```

```
In [60]: a=str((22+100)/(22+45+12+100))
print('Correct rate is '+ a +'percent')
```

Correct rate is 0.6815642458100558percent

## Random Forest

```
In [83]: model_RF = RandomForestClassifier(max_depth=3, random_state=20)
model_RF.fit( train_x, train_y)
pred_y=model_RF.predict(test_x)
acc_log_reg=round(model_RF.score(test_x, test_y)*100, 2)
print(str(acc_log_reg)+'percent')
```

83.24percent

```
In [80]: confusion_matrix(test_y, pred_y, labels=[1, 0])
```

```
Out[80]: array([[ 40,  27],
               [   3, 109]])
```

```
In [81]: a=str((40+109)/(40+27+3+109))
print('Correct rate is '+ a +'percent')
```

Correct rate is 0.8324022346368715percent

```
In [ ]: # I have trained 4 model to make the classification. There are Logistic Regre
# Descision tree, SVC, and Random Forest. At first, I think the Random Forest
# best correct rate. However, Logistic Regression has the same correct rate a
# forest model. As a result, I will use both model to run the Hold_out testin
# which model is the best.
```

## Step 3 Run the hand\_out testing

```
In [114... df1=pd.read_csv('holdout_test.csv')
```

```
In [115... df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
```

```
Data columns (total 12 columns):
#      Column      Non-Null Count  Dtype
---  -
0      Survived      0 non-null      float64
1      PassengerId    418 non-null    int64
2      Pclass          418 non-null    int64
3      Name            418 non-null    object
4      Sex             418 non-null    object
5      Age            332 non-null    float64
6      SibSp           418 non-null    int64
7      Parch           418 non-null    int64
8      Ticket          418 non-null    object
9      Fare            417 non-null    float64
10     Cabin           91 non-null     object
11     Embarked        418 non-null    object
dtypes: float64(3), int64(4), object(5)
memory usage: 39.3+ KB
```

```
In [120... df1=df1.drop(columns=['PassengerId','Name','Cabin','Ticket'])
df1['family']=df1['SibSp']+df1['Parch']
df1=df1.drop(columns=['SibSp','Parch'])
df1['is_alone']=df1['family'].apply(lambda x:1 if x==0 else 0)
df1['is_female']= df1['Sex'].apply(lambda x:1 if x =='female' else 0 )
df1=df1.drop(columns=['Sex'])
df1['Age']=df1['Age'].interpolate()
df1['is_child']=df1['Age'].apply(lambda x:1 if x<18 else 0)
df1['Fare']=df1['Fare'].fillna(mean(df1[df1['Fare'].notnull()][ 'Fare']))
df1=df1.drop(columns=['Survived'])
df1=pd.get_dummies(df1,prefix='is')
```

```
In [121... df1.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 418 entries, 0 to 417
Data columns (total 10 columns):
#      Column      Non-Null Count  Dtype
---  -
0      Pclass          418 non-null    int64
1      Age             418 non-null    float64
2      Fare            418 non-null    float64
3      family          418 non-null    int64
4      is_alone        418 non-null    int64
5      is_female       418 non-null    int64
6      is_child        418 non-null    int64
7      is_C            418 non-null    uint8
8      is_Q            418 non-null    uint8
9      is_S            418 non-null    uint8
dtypes: float64(2), int64(5), uint8(3)
memory usage: 24.2 KB
```

```
In [122... model_tree =DecisionTreeClassifier(max_depth=3, random_state=20)
model_tree.fit(train_x,train_y)
pred_df=model_tree.predict(df1)
```

```
In [125... df_help=pd.read_csv('holdout_test.csv')
pred_result=pd.DataFrame({
    'PassengerId':df_help['PassengerId'],
    'Survived':pred_df
})
```

```
In [128... pred_result.to_csv('Titanic_result_ccwang4.csv',index=False)
```

In [ ]: